Enhancing Cardiovascular Disease Diagnosis: A Machine Learning and Ensemble Pipeline Using the UCI Heart Disease Dataset

Abstract

Cardiovascular diseases (CVDs) continue to be the leading cause of morbidity and mortality worldwide, responsible for approximately 18 million deaths each year, representing 31% of global deaths. *(World Health Organization, 2025)* Early and accurate diagnosis of CVD is essential to improving patient outcomes and reducing healthcare costs associated with advanced interventions and hospitalisations. *(Kumar et al., 2025)* Despite the availability of traditional diagnostic techniques such as clinical examination, electrocardiography (ECG), and biochemical markers, diagnosis remains challenging due to the heterogenous presentation of cardiovascular pathology and the complexity of risk factors. *(Ahsan & Siddique, 2022) & (Teja & G Mokesh Rayalu, 2025)* Machine Learning (ML), an area of artificial intelligence (AI), provides a framework to analyse large and multidimensional clinical datasets, uncovering hidden patterns and supporting risk stratification in ways unachievable by classical methods. *(Kumar et al., 2025)* Importantly, advancements in ensemble ML models combine the strength of multiple classifiers to enhance predictive accuracy and robustness. *(Ganie et al., 2025)*

This research proposes the development and benchmarking of individual and ensemble ML algorithms using the well established UCI Cleveland Heart Disease dataset. The project integrates advanced data preprocessing stages including missing data imputation and feature standardisation, feature engineering such as principal component analysis (PCA) and recursive feature elimination. As well as ensemble learning strategies, comprising hard and soft voting classifiers, stacking, and boosting architectures. *(Zaidi et al., 2023) & (Ashika & Hannah Grace, 2025)*. Methodological validity will be ensured through extensive hyperparameter tuning, stratified k-fold cross-validation, and statistical tests of significance. *(Cai et al., 2024)* Emphasis will be placed on interpretability through explainable AI methods like SHAP and LIME, facilitating clinical adoption by assessing the contribution of individual predictors. *(Ganie et al., 2025)* The anticipated outcome is a highly accurate ML pipeline, exceeding 90% diagnostic accuracy, coupled with transparent insights into key clinical risk factors. *(Waleed Alsabhan & Abdullah Alfadhly, 2025)* This study aims to contribute to both academic knowledge, through reproducible benchmarking, methodological comparison and feature analysis, and clinically actionable tools that support earlier diagnosis and personalised cardiovascular care. *(Kumar et al., 2025)* & *(Ganie et al., 2025)*

Problem Statement

Cardiovascular disease is a multifaceted health condition that accounts for significant mortality. *(World Health Organization, 2025)* Diagnosing heart disease is impeded by diverse clinical presentations, nonlinear interactions among risk factors, and limitations associated to traditional diagnostic protocols including ECG presentations, blood biomarkers, and physician experience. *(Kumar et al., 2025)* & *(Teja & G Mokesh Rayalu, 2025)* Despite advances, these methods inadequately model the multivariate clinical realities in disease progression, contributing to delays in diagnosis, misdiagnosis, and poor patient management. *(Ahsan & Siddique, 2022)* The complexity of cardiovascular pathophysiology necessitates approaches capable of synthesising high-dimensional clinical and demographic variables effectively. *(Zaidi et al., 2023)*

Machine learning techniques have emerged as promising tools to address these diagnostic gaps by identifying complex associations within large datasets. *(Kumar et al., 2025) & (Waleed Alsabhan & Abdullah Alfadhly, 2025)*. Predictive models derived from ML are able to discern subtle nonlinear relationships and interactions between variables, potentially outperforming conventional statistical methods. Furthermore, the continuous growth of electronic health records (EHRs) and biomedical databases enables the application of ML over extensive, heterogenous datasets, promising more individualised, data-driven risk assessments. *(Ganie et al., 2025) & (Toma et al., 2024)*

Nevertheless, several challenges persist with the translation of ML innovations into clinical practice. This includes the risks of overfitting on limited or unrepresentative datasets, lack of transparency in "black box" algorithms affecting clinician trust, limited generalisability to diverse populations due to sparse external validation studies, and operational challenges integrating ML outputs into existing clinical decision workflows. *(Kumar et al., 2025)* & *(Cai et al., 2024)* Although ensemble an deep learning architectures demonstrate superior classification accuracies, often exceeding 90%, there remains a need to systematically validate these findings across independent datasets and to manage redundant or confounding features that may distort predictions. *(Ganie et al., 2025) & (Ashika & Hannah Grace, 2025)* Without addressing such gaps, the clinical utility of ML algorithms will remain constrained, underlining an urgent research need. *(Zaidi et al., 2023)*

Research Objectives

This research aims to comprehensively evaluate and benchmark a diverse group of ML models for predicting the presence of cardiovascular disease using the Cleveland dataset, curated from the UCI ML Repository. Specific objectives include:

1. Investigating the impact of feature engineering and dimensionality reduction methods (PCA, recursive feature elimination (RFE), and random forest feature importance rankings) on both predictive accuracy and interpretability to streamline models and facilitate clinical understanding. *(Kumar et al., 2025) & (Zaidi et al., 2023)*
2. Comparing the diagnostic performance of individual classifiers including logistic regression, support vector machines (SVM), random forests, k-nearest neighbours (KNN), decision tree models, gradient boosting machines, XGBoost, CatBoost, and

multilayer perceptron's, incorporating systematic hyper parameter tuning and stratified k-fold cross-validation to optimise model generalisability. *(Kumar et al., 2025) & (Waleed Alsabhan & Abdullah Alfadhly, 2025)*

3. Developing and evaluating ensemble learning strategies, including hard voting, soft voting, stacking ensembles, and boosting algorithms, to enhance robustness and predictive accuracy beyond single model solutions. *(Ganie et al., 2025) & (Ashika & Hannah Grace, 2025)*
4. Applying explainable AI methods such as SHAP and LIME, to generate actionable insights into feature contributions, promoting clinician trust and enabling transparent decision-making frameworks. *(Ganie et al., 2025)*
5. Identifying practical barriers and enablers for real-world clinical adoption of ML models, including model complexity, interpretability, computational efficiency, and integration pathways within existing healthcare environments. *(Toma et al., 2024)*

The methodological design of this study is guided by several core research questions. It seeks to determine which feature selection techniques, such as, filter, wrapper, and embedded methods including mutual information, RFE, and PCA, are most effective in optimising predictive performance and clinical relevance within heart disease models. *(Kumar et al., 2025) & (Khan et al., 2024)*. The research investigates which ensemble learning approaches, notably hard voting, soft voting, stacking, and boosting, proving the optimal balance between diagnostic accuracy and interpretability when validated thoroughly with cross-validation and statistical analysis. *(Ashika & Hannah Grace, 2025)* There is a focus on evaluating how interpretability frameworks, such as SHAP and LIME, can be systematically integrated into ML workflows to foster user confidence and facilitate adoption among clinicians and healthcare organisations. *(Ganie et al., 2025)* Finally, the study addresses practical implementation issues, examining what technical, organisational, and workflow considerations are essential for successfully translating ML models from experimental or development setting into impactful tools for routine clinical practice. *(Zaidi et al., 2023)* By leveraging systematic feature selection, robust dimensionality reduction, and explainable frameworks, this work seeks to answer the question:

Can a validated, interpretable, and clinically integrated ensemble machine learning pipeline meaningfully outperform current standard diagnostic approaches for cardiovascular disease, enabling earlier intervention and improved patient outcomes in clinical contexts?

Literature Review

Machine learning (ML) has greatly improved how CVD risk is predicted and diagnosed, helping doctors find hidden patterns in complex clinical data that traditional methods often miss. *Kumar et al. (2025)* provide a detailed overview of both traditional and newer deep learning methods and highlight how ensemble techniques are becoming more popular because they can improve accuracy and reliability. *(Ganie et al., 2025)* These ensemble models reduce errors that can happen with single models and provide consistent predictions.

Some of the best performing algorithms for heart disease diagnosis include tree-based models like Random Forests, Gradient Boosting Machines (GBM), and XGBoost. Studies show these models often achieve over 90% accuracy using well known datasets such as the UCI Cleveland dataset. *(Ahsan & Siddique, 2022)* & *(Teja & G Mokesh Rayalu, 2025)* These models work well because they can handle complex nonlinear relationships and interactions among clinical features that simpler methods may miss. For example *Ahsan and Siddique (2022)* showed that XGBoost does better than many other algorithms when dealing with mixed data and missing values.

Managing the large number of features in cardiovascular datasets is important for improving model performance. Techniques like PCA, which reduce dataset size by combining related features, and feature selection methods like recursive feature elimination (RFE), play a key role in simplifying models and making them easier to understand. *(Zaidi et al., 2023)* These techniques help highlight key clinical features, such as, patient age, chest pain type, maximum heart rate, and exercise-induced angina, that contribute most to accurate predictions. By simplifying the data, these methods also help overfitting, so models perform better on new data.

Another big improvement in ML for heart disease has come from Explainable AI (XAI). Many ML models are "black boxes" that do not clearly show how decisions are made. Tools like SHAP and LIME make these models more transparent by explaining the contribution of each feature to a prediction. *(Ganie et al., 2025)* & *(Ashika & Hannah Grace, 2025)*. This helps build trust among clinicians, who need to understand why a model made a certain diagnosis before using it in practice. Some advanced models combine boosting and stacking approaches to balance high accuracy with clear interpretability, making them especially useful in clinical settings.

Despite this progress, challenges still slow the use of ML in regular CVD diagnosis. Many models are tested only in single hospitals or small patient groups, which limits how well they work on different populations. *(Cai et al., 2024)* Clinical data can be very diverse, and patient privacy rules create challenges for sharing data and models widely. *(Kumar et al., 2025) & (Toma et al., 2024)* Integrating ML tools into busy clinical workflows without causing disruptions is also difficult. To face these issues, researchers are exploring federated machine learning, which allows model training on data stored at multiple locations without sharing patient information, protecting privacy while improving performance. *(Cai et al., 2024)* Open datasets from multiple hospitals and shared benchmarking platforms also help create standards and promote replicable research. *(Toma et al., 2024)*

Multimodal data integration is another growing trend in improving CVD prediction. Studies have shown that combining different types of patient data, and wearables can improve early detection and risk stratification. For instance, integrating heart imaging data or genetic risk scores has improved diagnostic accuracy and helped identify people at higher risk earlier. *(Teja & G Mokesh Rayalu, 2025)* While this approach creates new challenges for combining and analysing complex, heterogenous data, it also offers promising paths to more personalised, precise care.

Recent advancements also highlight the impact of digital health tools such as wearable health devices and mobile applications that continuously collect heart rate and physical activity data. When combined with ML algorithms, these tools enable real-time monitoring and early warning systems for cardiovascular events, marking a shift toward proactive and preventative healthcare. These technologies help expand the reach of ML-based diagnostics beyond traditional clinical environments into everyday patient care.

In summary, ML methods for predicting heart disease have greatly improved, with ensemble models, feature reduction techniques, and XAI playing major roles. However, more work is needed to ensure these models work well across different patient groups, protect privacy, comply with regulations, and fit into clinical practice seamlessly. This evolving research field continues to move toward tools that are not only technically accurate, but also practical, interpretable and accessible for widespread clinical use.

Methodology

This research utilises the UCI Cleveland Heart Disease dataset, comprising 303 records with 14 clinical and demographic features including both continuous and categorical variables measured at standard clinical intervals. Data features include age, sex, chest pain type, resting blood pressure, serum cholesterol, fasting blood sugar, resting ECG results, maximal heart rate, exercise-induced angina, ST depression induced fluoroscopy, and thalassemia type.

The preprocessing pipeline involves data cleaning, beginning with missing data imputation through KNN, selected for effectiveness with mixed data types and minimal bias introduction. *(Cai et al., 2024)* Categorical variables are encoded with one hot encoding to avoid collinearity and support most ML algorithms. Normative scaling of numeric features via z-score standardisation ensures uniform feature contribution in distance and gradient-based models, optimising convergence and performance.

Feature engineering employs recursive feature elimination combined with random forest importance rankings to reduce dimensionality and enhance model interpretability. *(Zaidi et al., 2023)* PCA is applied for covariance reduction, retaining principal components explaining 85-90% of variance, allowing simplification without loosing information. Complementary nonlinear dimensionality reduction techniques such as t-SNE and UMAP visualise structure and patient clusters, providing insights into data separability and guiding feature combinations.

The study implements a suite of classifiers: logistic regression for baseline linear modelling, kernel-based SVM, KNN, decision trees, random forest as ensemble bagging approach, gradient boosting, and XGBoost and CatBoost algorithms, alongside multilayer perceptron (MLP) neural networks for nonlinear classification. *(Kumar et al., 2025) & (Ahsan & Siddique, 2022)* Hyperparameter tuning via grid and random search ensures optimal model configurations. Stratified k-fold validation is used to reduce variance in evaluation metric and provide realistic estimates based on sample performance. *(Teja & G Mokesh Rayalu, 2025)*

Ensemble methods constructed include hard voting based on majority vote, soft voting aggregating probabilistic outputs, and meta-learning stacking classifiers integrating multiple base classifiers through a second stage learner . *(Ganie et al., 2025)* Model outputs are further analysed with explainable AI techniques such as SHAP and LIME to provide global and local interpretability, essential for clinician acceptance. *(Ashika & Hannah Grace, 2025)*

Model evaluations rely on comprehensive metrics: accuracy, precision, recall, F1 score, area under the receiver operating characteristic curve (ROC-AUC), precision-recall AUC (PR-AUC), and calibration measures. Statistical significance in differences between models is assessed with McNemar's test and bootstrapped confidence intervals, providing comparative analysis. *(Cai et al., 2024)* All computational procedures utilise Python libraries including scikit-learn, XGBoost, CatBoost, pandas, SHAP, and visualisation modules, with an emphasis on full transparency and reproducibility through publicly available code repositories.

Limitations acknowledged include the retrospective and single source nature of the dataset, acknowledged class imbalance requiring synthetic oversampling, and potential lack of generalisability on real populations, which will be mitigated by validation on test data and model calibration. *(Cai et al., 2024)*

Expected Results

The project expects to demonstrate a high-performing ensemble ML pipeline for cardiovascular disease diagnosis, achieving predictive accuracies and ROC-AUC values exceeding 90%, exceeding individual classifier benchmarks on the UCI Cleveland dataset. *(Kumar et al., 2025)* Explainable AI support asserts the relative importance of clinical features including chest pain characteristics, maximal heart rate, serum cholesterol levels, and the extent of major vessel involvement, giving clearer insight for clinical decision making. *(Ganie et al., 2025) & (Ashika & Hannah Grace, 2025)*

Further, the research will contribute valuable benchmarks for comparing numerous different model architectures and feature engineering approaches in cardiovascular diagnostics. By evaluating hyperparameter tuned classifiers and ensembles with multiple validation metrics, this work will identify optimal combinations that can be used as best practice templates for future ML development in cardiology. *(Waleed Alsabhan & Abdullah Alfadhly, 2025)* The inclusion of balanced datasets through SMOTE and thorough calibration will ensure robustness and generalisability, mitigating common issues such as overfitting and class imbalance. Additionally, the reproducibility of findings and open availability of code and methodologies will provide foundation for further scientific inquiry and clinical translation, with an aim of wider adoption. These outcomes align with the broader healthcare goals of cost effective and early cardiovascular disease detection, a major factor in reducing mortality and improving quality of life globally.

Timeline

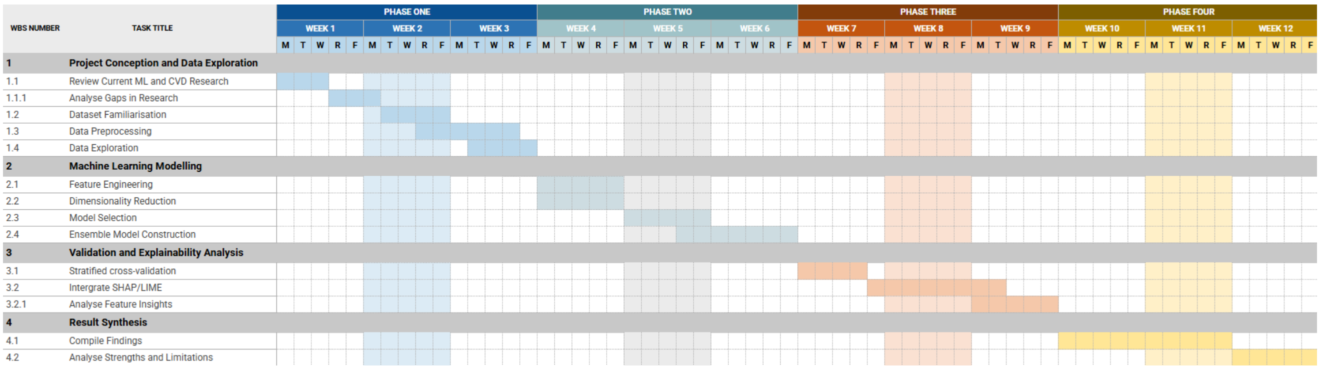

References

Ahsan, M. M., & Siddique, Z. (2022). Machine learning-based Heart Disease diagnosis: a Systematic Literature Review. *Artificial Intelligence in Medicine*, *128*, 102289. https://doi.org/10.1016/j.artmed.2022.102289

Alqahtani, A., Alsubai, S., Sha, M., Vilcekova, L., & Javed, T. (2022). Cardiovascular Disease Detection using Ensemble Learning. *Computational Intelligence and Neuroscience*, *2022*, 1–9. https://doi.org/10.1155/2022/5267498

Arian Zargarzadeh, Elnaz Javanshir, Ghaffari, A., Erfan Mosharkesh, & Babak Anari. (2023). Artificial intelligence in cardiovascular medicine: An updated review of the literature. *Journal of Cardiovascular and Thoracic Research. (Print)*, *15*(4), 204–209. https://doi.org/10.34172/jcvtr.2023.33031

Ashika, T., & Hannah Grace, G. (2025). Enhancing heart disease prediction with stacked ensemble and MCDM-based ranking: an optimized RST-ML approach. *Frontiers in Digital Health*, *7*. https://doi.org/10.3389/fdgth.2025.1609308

Bayona, A. V., Wang, J., & Xiang, Y. (2025). Artificial intelligence in cardiovascular prognosis and diagnosis: a review. *Exploration of Medicine*, *6*. https://doi.org/10.37349/emed.2025.1001347

Cai, Y.-Q., Cai, Y., Tang, L.-Y., Jing, T.-C., Gong, M., Li, H.-J., Hu, W., Zhang, X., Gong, D.-X., & Zhang, G.-W. (2024). Pitfalls in Developing Machine Learning Models for Predicting Cardiovascular Diseases: Challenge and Solutions (Preprint). *Journal of Medical Internet Research*, *26*, e47645–e47645. https://doi.org/10.2196/47645

Ganie, S. M., Pramanik, P. K. D., & Zhao, Z. (2025). Ensemble learning with explainable AI for improved heart disease prediction based on multiple datasets. *Scientific Reports*, *15*(1). https://doi.org/10.1038/s41598-025-97547-6

Hamid, M., Fahima Hajjej, Ala Saleh Alluhaidan, & Waleed, N. (2025). Fine tuned CatBoost machine learning approach for early detection of cardiovascular disease through predictive modeling. *Scientific Reports*, *15*(1). https://doi.org/10.1038/s41598-025-13790-x

Khan, M. A., Mazhar, T., Mateen Yaqoob, M., Badruddin Khan, M., Jilani Saudagar, A. K., Ghadi, Y. Y., Khattak, U. F., & Shahid, M. (2024). Optimal feature selection for heart disease prediction using modified Artificial Bee colony (M-ABC) and K-nearest neighbors (KNN). *Scientific Reports*, *14*(1). https://doi.org/10.1038/s41598-024-78021-1

Kumar, R., Garg, S., Kaur, R., Johar, M., Singh, S., Menon, S. V., Kumar, P., Hadi, A. M., Hasson, S. A., & Jasmina Lozanović. (2025). A comprehensive review of machine learning for heart disease prediction: challenges, trends, ethical considerations, and future directions. *Frontiers in Artificial Intelligence*, *8*. https://doi.org/10.3389/frai.2025.1583459

Teja, M. D., & G Mokesh Rayalu. (2025). Optimizing heart disease diagnosis with advanced machine learning models: a comparison of predictive performance. *BMC Cardiovascular Disorders*, *25*(1). https://doi.org/10.1186/s12872-025-04627-6

Toma, M., Chong, L., Husain, G., Nasef, D., Vathappallil, P., & Mihir Matalia. (2024). Machine Learning Strategies for Improved Cardiovascular Disease Detection. *Medical Research Archives*, *13*(1). https://doi.org/10.18103/mra.v13i1.6245

Waleed Alsabhan, & Abdullah Alfadhly. (2025). Effectiveness of machine learning models in diagnosis of heart disease: A comparative study. *Scientific Reports*, *15*(1). https://doi.org/10.1038/s41598-025-09423-y

World Health Organization. (2025, July 31). *Cardiovascular diseases (CVDs)*. World Health Organization . https://www.who.int/news-room/fact-sheets/detail/cardiovascular-diseases-(cvds)

Zaidi, H. A., Jones, R. E., Hammersley, D. J., Hatipoglu, S., Balaban, G., Mach, L., Halliday, B. P., Lamata, P., Prasad, S. K., & Bishop, M. J. (2023). Machine learning analysis of complex late gadolinium enhancement patterns to improve risk prediction of major arrhythmic events. *Frontiers in Cardiovascular Medicine*, *10*. https://doi.org/10.3389/fcvm.2023.1082778